

```

# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

/kaggle/input/coc-npn/customer_support_tickets.csv
/kaggle/input/coc-npn/WA_Fn-UseC_-Telco-Customer-Churn.csv
/kaggle/input/coc-npn/CDR-Call-Details.csv

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
plt.style.use('ggplot')
# pd.set_option('max_columns', 100)

```

EDA

<https://www.kaggle.com/code/shreeramt/churn-upsell/>

```
df = pd.read_csv("/kaggle/input/coc-npn/CDR-Call-Details.csv")
```

understanding the data

```
df.shape
```

```
(101174, 17)
```

```
df.columns
Index(['Phone Number', 'Account Length', 'VMail Message', 'Day Mins',
      'Day Calls', 'Day Charge', 'Eve Mins', 'Eve Calls', 'Eve
Charge',
      'Night Mins', 'Night Calls', 'Night Charge', 'Intl Mins', 'Intl
Calls',
      'Intl Charge', 'CustServ Calls', 'Churn'],
      dtype='object')
```

```
df = df.rename(columns={
    'Phone Number': 'Phone_Number',
    'Account Length': 'Account_Length',
    'VMail Message': 'VMail_Message',
    'Day Mins': 'Day_Mins',
    'Day Calls': 'Day_Calls',
    'Day Charge': 'Day_Charge',
    'Eve Mins': 'Eve_Mins',
    'Eve Calls': 'Eve_Calls',
    'Eve Charge': 'Eve_Charge',
    'Night Mins': 'Night_Mins',
    'Night Calls': 'Night_Calls',
    'Night Charge': 'Night_Charge',
    'Intl Mins': 'Intl_Mins',
    'Intl Calls': 'Intl_Calls',
    'Intl Charge': 'Intl_Charge',
    'CustServ Calls': 'CustServ_Calls'
})
```

```
df
```

	Phone_Number	Account_Length	VMail_Message	Day_Mins
Day_Calls \				
0	382-4657	128	25	265.1
110				
1	371-7191	107	26	161.6
123				
2	358-1921	137	0	243.4
114				
3	375-9999	84	0	299.4
71				
4	330-6626	75	0	166.7
113				
...
.				
101169	789-9756	222	0	228.2
60				
101170	798-5885	88	0	282.2
222				
101171	798-5798	22	0	222.2

62					
101172	999-9897	228	0	222.0	
99					
101173	786-7589	228	0	226.2	
98					

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge	Night_Mins
Night_Calls \					
0	45.07	197.4	99	16.78	244.7
91					
1	27.47	195.5	103	16.62	254.4
103					
2	41.38	121.2	110	10.30	162.6
104					
3	50.90	61.9	88	5.26	196.9
89					
4	28.34	148.3	122	12.61	186.9
121					
...
...					
101169	22.82	229.8	289	28.26	222.8
222					
101170	82.88	208.8	220	22.82	282.2
200					
101171	88.66	228.0	228	22.08	62.2
209					
101172	88.08	220.2	80	22.92	282.9
28					
101173	86.28	288.2	208	28.28	800.0
228					

	Night_Charge	Intl_Mins	Intl_Calls	Intl_Charge
CustServ_Calls \				
0	11.01	10.0	3	2.70
1				
1	11.45	13.7	3	3.70
1				
2	7.32	12.2	5	3.29
0				
3	8.86	6.6	7	1.78
2				
4	8.41	10.1	3	2.73
3				
...
..				
101169	2.28	6.2	2	2.62
2				
101170	20.68	9.8	8	2.82
8				

```

101171      2.26      2.8      6      2.22
2
101172      20.22     2.2      8      0.82
0
101173      28.80     20.0      8      2.20
2

```

```

      Churn
0      False
1      False
2      False
3      False
4      False
...
101169  False
101170  False
101171  False
101172  False
101173  False

```

```
[101174 rows x 17 columns]
```

```

# df['day_avg'] = df['Day Mins'] / df['Day Calls']
# df['day_avg']

```

```
df.head()
```

	Phone_Number	Account_Length	VMail_Message	Day_Mins	Day_Calls	\
0	382-4657	128	25	265.1	110	
1	371-7191	107	26	161.6	123	
2	358-1921	137	0	243.4	114	
3	375-9999	84	0	299.4	71	
4	330-6626	75	0	166.7	113	

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge	Night_Mins	Night_Calls \
0	45.07	197.4	99	16.78	244.7	
91						
1	27.47	195.5	103	16.62	254.4	
103						
2	41.38	121.2	110	10.30	162.6	
104						
3	50.90	61.9	88	5.26	196.9	
89						
4	28.34	148.3	122	12.61	186.9	
121						

	Night_Charge	Intl_Mins	Intl_Calls	Intl_Charge	CustServ_Calls	Churn
0	11.01	10.0	3	2.70		1

```
False
1      11.45      13.7      3      3.70      1
False
2      7.32      12.2      5      3.29      0
False
3      8.86      6.6      7      1.78      2
False
4      8.41      10.1      3      2.73      3
False
```

```
df.tail()
```

```
      Phone_Number  Account_Length  VMail_Message  Day_Mins
Day_Calls \
101169      789-9756           222           0      228.2
60
101170      798-5885           88           0      282.2
222
101171      798-5798           22           0      222.2
62
101172      999-9897          228           0      222.0
99
101173      786-7589          228           0      226.2
98
```

```
      Day_Charge  Eve_Mins  Eve_Calls  Eve_Charge  Night_Mins
Night_Calls \
101169      22.82      229.8      289      28.26      222.8
222
101170      82.88      208.8      220      22.82      282.2
200
101171      88.66      228.0      228      22.08      62.2
209
101172      88.08      220.2       80      22.92      282.9
28
101173      86.28      288.2      208      28.28      800.0
228
```

```
      Night_Charge  Intl_Mins  Intl_Calls  Intl_Charge
CustServ_Calls \
101169      2.28      6.2      2      2.62
2
101170      20.68      9.8      8      2.82
8
101171      2.26      2.8      6      2.22
2
101172      20.22      2.2      8      0.82
0
101173      28.80      20.0      8      2.20
2
```

```

Churn
101169 False
101170 False
101171 False
101172 False
101173 False

```

```
print(df.info()) #objects are strings essentially
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 101174 entries, 0 to 101173
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Phone_Number          101174 non-null object
1   Account_Length        101174 non-null int64
2   VMail_Message         101174 non-null int64
3   Day_Mins              101174 non-null float64
4   Day_Calls             101174 non-null int64
5   Day_Charge            101174 non-null float64
6   Eve_Mins              101174 non-null float64
7   Eve_Calls             101174 non-null int64
8   Eve_Charge            101174 non-null float64
9   Night_Mins            101174 non-null float64
10  Night_Calls           101174 non-null int64
11  Night_Charge          101174 non-null float64
12  Intl_Mins             101174 non-null float64
13  Intl_Calls            101174 non-null int64
14  Intl_Charge           101174 non-null float64
15  CustServ_Calls        101174 non-null int64
16  Churn                 101174 non-null bool
dtypes: bool(1), float64(8), int64(7), object(1)
memory usage: 12.4+ MB
None

```

```
print(df.describe())
```

	Account_Length	VMail_Message	Day_Mins	Day_Calls	\
count	24361.000000	24361.000000	24361.000000	24361.000000	
mean	142.557982	7.007143	235.520038	144.334017	
std	93.875565	12.810141	50.888264	83.959842	
min	1.000000	0.000000	62.300000	20.000000	
25%	44.000000	0.000000	222.200000	84.000000	
50%	109.000000	0.000000	229.800000	109.000000	
75%	223.000000	0.000000	266.900000	222.000000	
max	329.000000	51.000000	444.400000	329.000000	

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge	\
Night_Mins					

count	24361.000000	24361.000000	24361.000000	24361.000000
24361.000000				
mean	33.434129	245.841118	145.180370	24.016874
244.790738				
std	16.377482	35.953885	83.943032	3.877371
37.154832				
min	6.220000	154.800000	12.000000	13.160000
131.600000				
25%	22.960000	222.400000	82.000000	22.220000
222.300000				
50%	28.820000	232.900000	113.000000	22.980000
232.300000				
75%	34.410000	280.200000	222.000000	26.620000
280.800000				
max	82.980000	400.900000	329.000000	32.990000
367.700000				

	Night_Calls	Night_Charge	Intl_Mins	Intl_Calls
Intl_Charge \				
count	24361.000000	24361.000000	24361.000000	24361.000000
24361.000000				
mean	143.544559	11.437698	13.342622	3.904109
2.709633				
std	83.582519	8.018891	8.078907	2.395398
0.722454				
min	20.000000	2.000000	0.000000	0.000000
0.000000				
25%	82.000000	4.240000	6.600000	2.000000
2.240000				
50%	108.000000	9.220000	9.800000	3.000000
2.490000				
75%	222.000000	20.290000	22.000000	6.000000
2.920000				
max	329.000000	32.980000	32.900000	16.000000
5.400000				

	CustServ_Calls
count	24361.000000
mean	2.046673
std	1.760284
min	0.000000
25%	2.000000
50%	2.000000
75%	2.000000
max	11.000000

dropping the churn column

From what I've seen `account_length` seems to indicate the duration the account was open so given the stat summary of it I assume it is measured in days which validates the presence of other features

```
# df = df.drop(['Churn'], axis=1).copy()
```

Null & NaN

```
print(df.isna().sum())
```

```
Phone_Number      0
Account_Length    0
VMail_Message     0
Day_Mins          0
Day_Calls         0
Day_Charge        0
Eve_Mins          0
Eve_Calls         0
Eve_Charge        0
Night_Mins        0
Night_Calls       0
Night_Charge      0
Intl_Mins         0
Intl_Calls        0
Intl_Charge       0
CustServ_Calls    0
Churn             0
dtype: int64
```

Duplicates

```
df.duplicated().sum()
# duplicated_rows = df[df.duplicated()]
# print(duplicated_rows)
```

```
40729
```

```
df.loc[df.duplicated()]
```

	Phone_Number	Account_Length	VMail_Message	Day_Mins	Day_Calls
9999	785-9657	228	28	268.2	220
10000	779-7999	202	26	262.6	228
10001	758-9959	282	0	228.2	222

10002	775-9999	82	0	299.2	22
10003	778-6656	28	0	266.2	228
...
99985	775-9988	292	86	286.2	22
99986	757-7869	68	0	282.2	82
99987	797-5759	28	0	280.8	209
99988	765-9799	282	0	228.8	208
99989	786-8989	22	28	282.2	228
<div> <div>Day_Charge</div> <div>Eve_Mins</div> <div>Eve_Calls</div> <div>Eve_Charge</div> <div>Night_Mins</div> <div>Night_Calls \</div> </div>					
9999	28.02	292.2	99	26.28	222.2
92					
10000	22.22	298.8	208	26.62	282.2
208					
10001	22.88	222.2	220	20.80	262.6
202					
10002	80.90	62.9	88	8.26	296.9
89					
10003	28.82	228.8	222	22.62	286.9
222					
...
...					
99985	26.88	228.8	226	28.82	229.2
88					
99986	89.29	288.2	88	28.02	292.8
228					
99987	80.22	288.8	88	22.88	292.9
92					
99988	86.88	289.6	82	28.82	289.2
282					
99989	89.88	268.9	82	22.60	222.2
22					
<div> <div>Night_Charge</div> <div>Intl_Mins</div> <div>Intl_Calls</div> <div>Intl_Charge</div> <div>CustServ_Calls</div> <div>Churn</div> </div>					
9999	22.02	20.0	8	2.20	
2 False					
10000	22.28	28.2	8	8.20	
2 False					
10001	2.82	22.2	8	8.29	
0 False					

10002	8.86	6.6	2	2.28
2 False				
10003	8.22	20.2	8	2.28
8 False				
...
...				..
99985	22.86	9.9	6	2.62
2 False				
99986	8.62	9.6	2	2.89
8 False				
99987	8.62	22.2	6	8.82
2 False				
99988	6.26	8.0	20	2.88
2 False				
99989	20.86	28.2	2	8.20
0 False				

[40729 rows x 17 columns]

df.query('Phone_Number == "779-7999"').head()

	Phone_Number	Account_Length	VMail_Message	Day_Mins	Day_Calls
\					
3334	779-7999	407	46	464.6	445
6667	779-7999	202	26	262.6	228
10000	779-7999	202	26	262.6	228
13333	779-7999	202	26	262.6	228
16666	779-7999	202	26	262.6	228

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge	Night_Mins
Night_Calls \					
3334	47.47	495.5	405	46.64	454.4
405					
6667	22.22	298.8	208	26.62	282.2
208					
10000	22.22	298.8	208	26.62	282.2
208					
13333	22.22	298.8	208	26.62	282.2
208					
16666	22.22	298.8	208	26.62	282.2
208					

	Night_Charge	Intl_Mins	Intl_Calls	Intl_Charge
CustServ_Calls Churn				
3334	44.45	45.7	5	5.7

```

4  False
6667      22.28      28.2      8      8.2
2  False
10000     22.28      28.2      8      8.2
2  False
13333     22.28      28.2      8      8.2
2  False
16666     22.28      28.2      8      8.2
2  False

```

```

df = df.loc[~df.duplicated()].reset_index(drop=True).copy()
df

```

	Phone_Number	Account_Length	VMail_Message	Day_Mins	Day_Calls
0	382-4657	128	25	265.1	110
1	371-7191	107	26	161.6	123
2	358-1921	137	0	243.4	114
3	375-9999	84	0	299.4	71
4	330-6626	75	0	166.7	113
...
60440	789-9756	222	0	228.2	60
60441	798-5885	88	0	282.2	222
60442	798-5798	22	0	222.2	62
60443	999-9897	228	0	222.0	99
60444	786-7589	228	0	226.2	98

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge	Night_Mins
0	45.07	197.4	99	16.78	244.7
1	27.47	195.5	103	16.62	254.4
2	41.38	121.2	110	10.30	162.6
3	50.90	61.9	88	5.26	196.9
4	28.34	148.3	122	12.61	186.9
...

```

...
60440      22.82      229.8      289      28.26      222.8
222
60441      82.88      208.8      220      22.82      282.2
200
60442      88.66      228.0      228      22.08      62.2
209
60443      88.08      220.2      80      22.92      282.9
28
60444      86.28      288.2      208      28.28      800.0
228

```

```

      Night_Charge Intl_Mins Intl_Calls Intl_Charge
CustServ_Calls  Churn
0      11.01      10.0      3      2.70
1  False
1      11.45      13.7      3      3.70
1  False
2      7.32      12.2      5      3.29
0  False
3      8.86      6.6      7      1.78
2  False
4      8.41      10.1      3      2.73
3  False

```

```

...      ...      ...      ...      ...      ...
.      ...
60440      2.28      6.2      2      2.62
2  False
60441      20.68      9.8      8      2.82
8  False
60442      2.26      2.8      6      2.22
2  False
60443      20.22      2.2      8      0.82
0  False
60444      28.80      20.0      8      2.20
2  False

```

[60445 rows x 17 columns]

unique counts in a feature

```
df.nunique()
```

```

Phone_Number      7467
Account_Length    322
VMail_Message      72
Day_Mins          2548
Day_Calls         221
Day_Charge        2873
Eve_Mins          2523

```

```
Eve_Calls      224
Eve_Charge     2221
Night_Mins     2464
Night_Calls    218
Night_Charge   1470
Intl_Mins      267
Intl_Calls     39
Intl_Charge    339
CustServ_Calls 11
Churn          2
dtype: int64
```

```
df.shape
```

```
(60445, 17)
```

Outliers

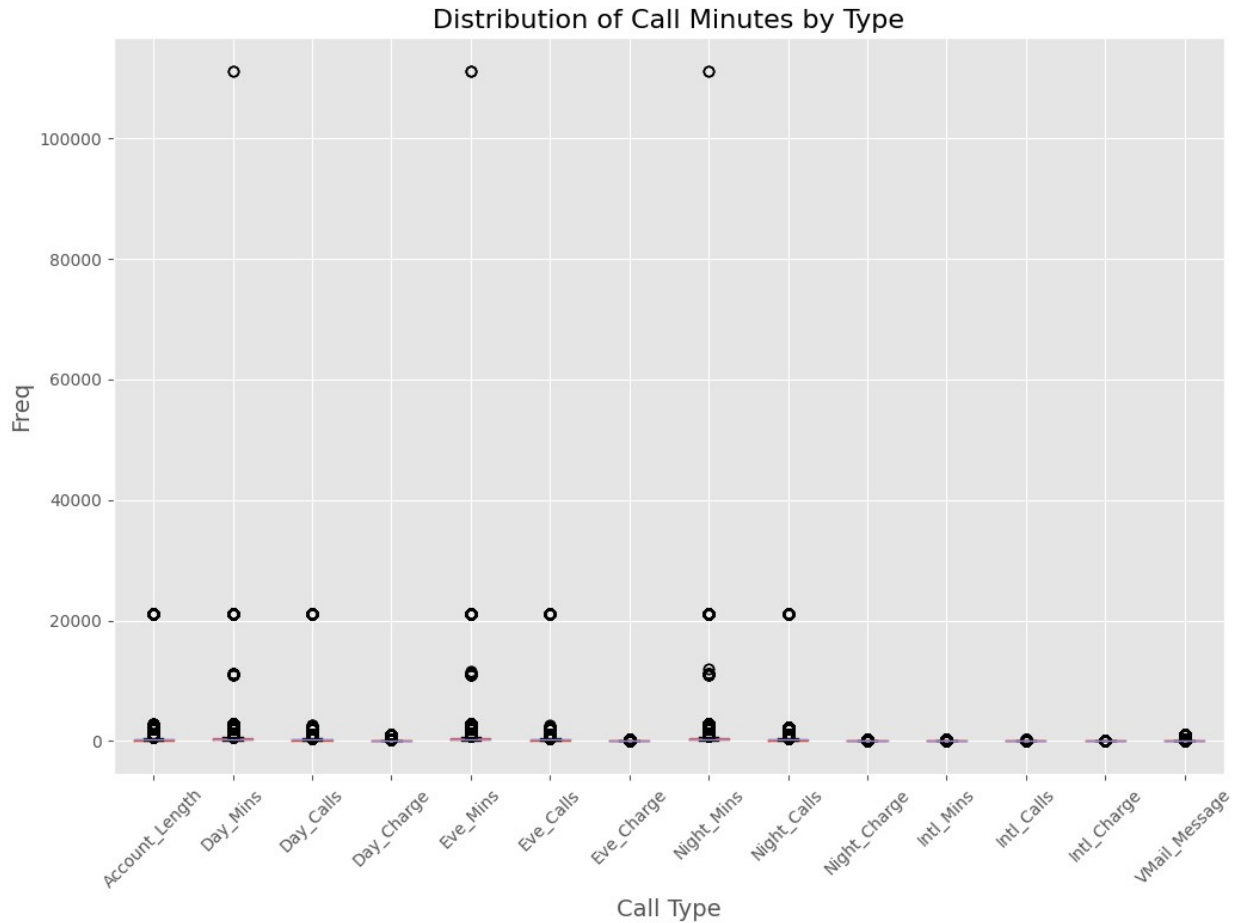
```
fig, ax = plt.subplots(figsize=(12, 8))

numeric_cols = ['Account_Length', 'Day_Mins', 'Day_Calls', 'Day_Charge',
                'Eve_Mins', 'Eve_Calls', 'Eve_Charge',
                'Night_Mins', 'Night_Calls', 'Night_Charge', 'Intl_Mins', 'Intl_Calls', 'In
tl_Charge', 'VMail_Message']

# Create a box plot
df.boxplot(column= numeric_cols,
           ax=ax)

# Add titles and labels
ax.set_title('Distribution of Call Minutes by Type', fontsize=16)
ax.set_xlabel('Call Type', fontsize=14)
ax.set_ylabel('Freq', fontsize=14)
ax.set_xticklabels(numeric_cols, rotation=45)

plt.show()
```



```
df.describe()
```

	Account_Length	VMail_Message	Day_Mins	Day_Calls
count	60445.000000	60445.000000	60445.000000	60445.000000
mean	329.385541	18.101613	596.555649	274.752138
std	1436.763064	76.859936	2231.437206	997.047766
min	1.000000	0.000000	0.000000	0.000000
25%	68.000000	0.000000	222.200000	88.000000
50%	200.000000	0.000000	262.200000	202.000000
75%	243.000000	22.000000	404.400000	226.000000
max	21111.000000	1111.000000	111111.110000	21111.000000

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge
Night_Mins				
count	60445.000000	60445.000000	60445.000000	60445.000000
mean	64.880486	669.288486	272.647911	34.894168
std	134.695269	2476.768325	956.454815	36.584588
min	0.000000	0.000000	0.000000	0.000000
max	20.300000			

25%	24.220000	223.400000	88.000000	22.220000
223.200000				
50%	32.920000	268.800000	200.000000	24.900000
269.110000				
75%	80.480000	440.400000	226.000000	32.620000
440.400000				
max	1111.990000	111111.200000	21111.000000	211.990000
111111.110000				

	Night_Calls	Night_Charge	Intl_Mins	Intl_Calls
Intl_Charge \				
count	60445.000000	60445.000000	60445.000000	60445.000000
60445.000000				
mean	267.133377	14.066812	20.571371	5.728894
4.315376				
std	913.240573	16.532483	25.494288	8.271802
2.711543				
min	20.000000	1.040000	0.000000	0.000000
0.000000				
25%	88.000000	6.220000	8.800000	2.000000
2.280000				
50%	200.000000	9.220000	20.200000	4.000000
2.920000				
75%	226.000000	20.400000	22.800000	8.000000
4.940000				
max	21111.000000	211.920000	211.900000	211.000000
11.920000				

	CustServ_Calls
count	60445.000000
mean	2.563438
std	2.376449
min	0.000000
25%	2.000000
50%	2.000000
75%	4.000000
max	11.000000

```
# df = df[
#     (df['Account_Length'] <= 20000) &
#     (df['Day_Mins'] <= 20000) &
#     (df['Eve_Mins'] <= 20000) &
#     (df['Night_Mins'] <= 20000) &
#     (df['Intl_Mins'] <= 20000) &
#     (df['Day_Calls'] <= 20000) &
#     (df['Eve_Calls'] <= 20000) &
#     (df['Night_Calls'] <= 20000) &
#     (df['Intl_Calls'] <= 20000) &
```

```

# (df['CustServ_Calls'] <= 20000)
# ].copy()

def remove_outliers_iqr(df, column_name):
    Q1 = df[column_name].quantile(0.25)
    Q3 = df[column_name].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    return df[(df[column_name] >= lower_bound) & (df[column_name] <=
upper_bound)]

# Loop through each numerical column
for col in numeric_cols:
    df = remove_outliers_iqr(df, col)

# Reset index after filtering
df.reset_index(drop=True, inplace=True)

df.shape

(24361, 17)


fig, ax = plt.subplots(figsize=(12, 8))

# box plot
df.boxplot(column=numeric_cols,
            ax=ax)

# titles and labels
ax.set_title('Distribution of Call Minutes by Type', fontsize=16)
ax.set_xlabel('Call Type', fontsize=14)
ax.set_ylabel('Freq', fontsize=14)
ax.set_xticklabels(numeric_cols, rotation=45)

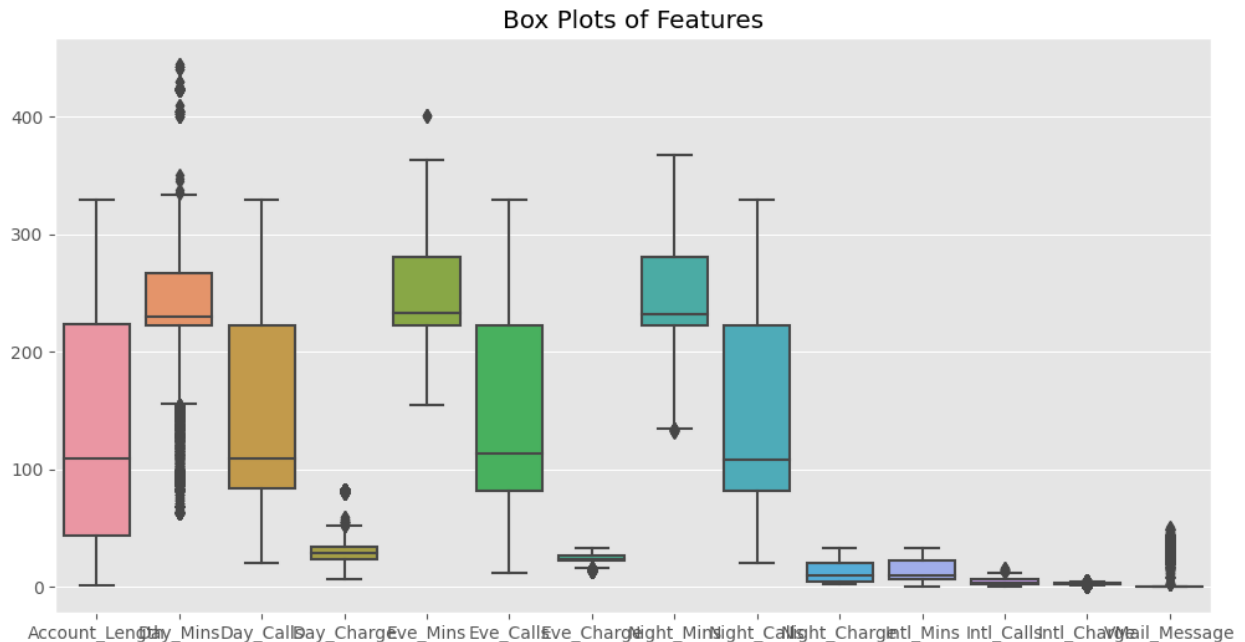
plt.show()

```

other visualizations

```
# boxplot
plt.figure(figsize=(12,6))
sns.boxplot(data=df[numeric_cols])
plt.title('Box Plots of Features')
plt.show()
```



```
df.describe()
```

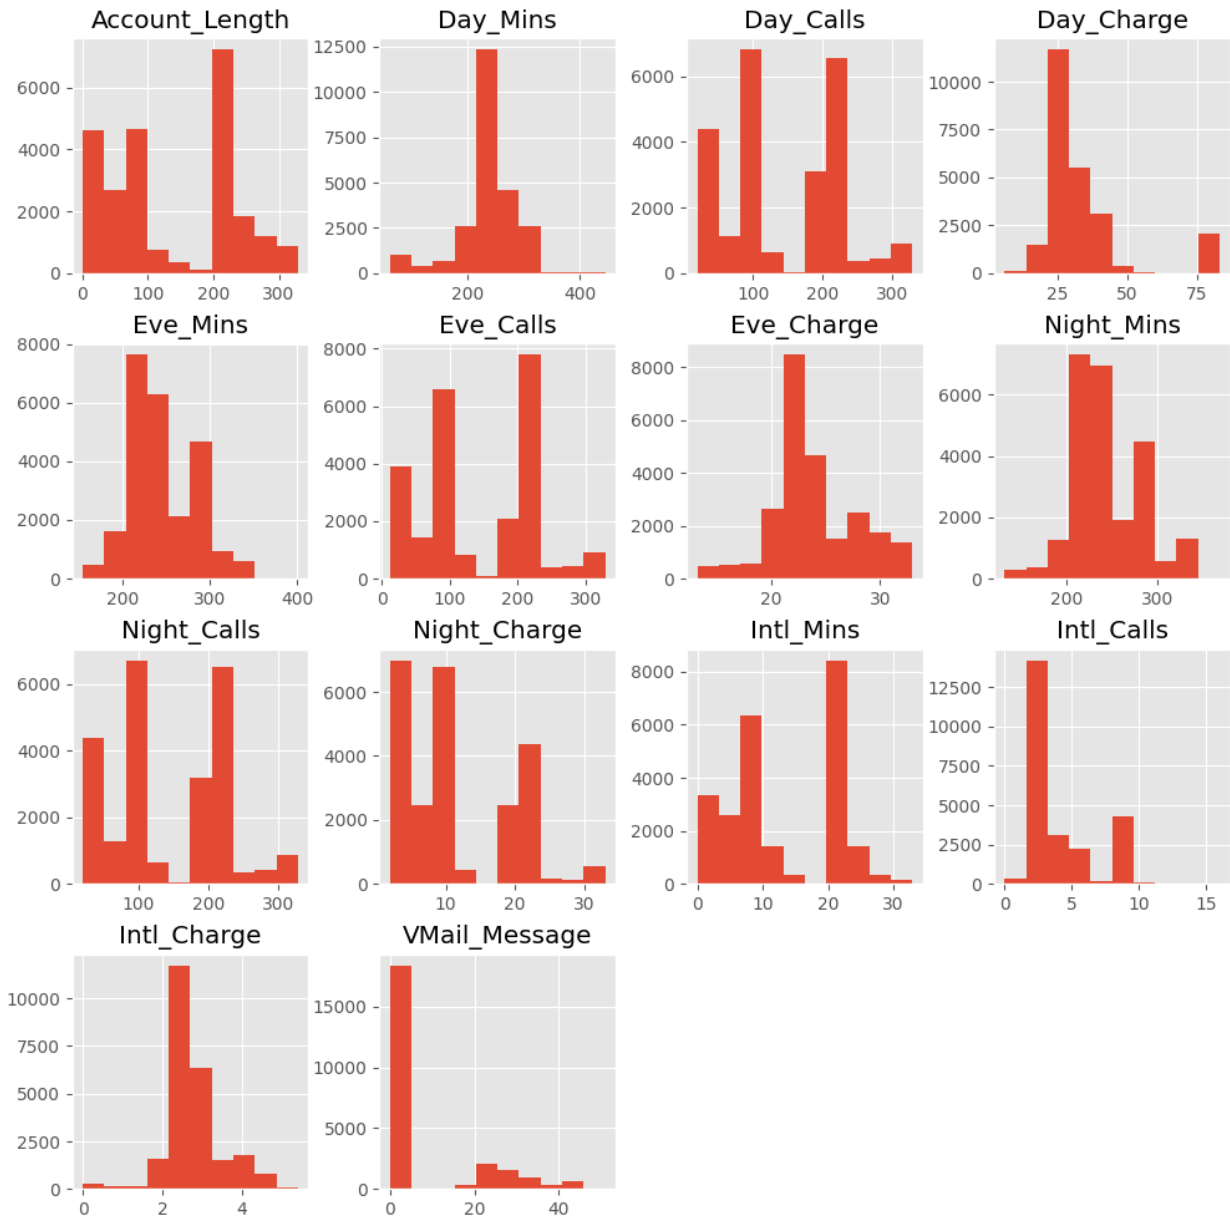
	Account_Length	VMail_Message	Day_Mins	Day_Calls
count	24361.000000	24361.000000	24361.000000	24361.000000
mean	142.557982	7.007143	235.520038	144.334017
std	93.875565	12.810141	50.888264	83.959842
min	1.000000	0.000000	62.300000	20.000000
25%	44.000000	0.000000	222.200000	84.000000
50%	109.000000	0.000000	229.800000	109.000000
75%	223.000000	0.000000	266.900000	222.000000
max	329.000000	51.000000	444.400000	329.000000

	Day_Charge	Eve_Mins	Eve_Calls	Eve_Charge
count	24361.000000	24361.000000	24361.000000	24361.000000
mean	33.434129	245.841118	145.180370	24.016874
std	16.377482	35.953885	83.943032	3.877371
min	6.220000	154.800000	12.000000	13.160000
25%	22.960000	222.400000	82.000000	22.220000
50%	28.820000	232.900000	113.000000	22.980000
75%	34.410000	280.200000	222.000000	26.620000
max	82.980000	400.900000	329.000000	32.990000


	Night_Calls	Night_Charge	Intl_Mins	Intl_Calls
Intl_Charge \				
count	24361.000000	24361.000000	24361.000000	24361.000000
mean	143.544559	11.437698	13.342622	3.904109
std	83.582519	8.018891	8.078907	2.395398
min	20.000000	2.000000	0.000000	0.000000
25%	82.000000	4.240000	6.600000	2.000000
50%	108.000000	9.220000	9.800000	3.000000
75%	222.000000	20.290000	22.000000	6.000000
max	329.000000	32.980000	32.900000	16.000000

	CustServ_Calls
count	24361.000000
mean	2.046673
std	1.760284
min	0.000000
25%	2.000000
50%	2.000000
75%	2.000000
max	11.000000

```
# histograms
df[numeric_cols].hist(figsize=(12, 12))
plt.show()
```




```
# correlation heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(df[numeric_cols].corr(), annot=True, fmt='.2f',
            cmap='magma')
plt.title('Correlation Heatmap')
plt.show()
```




```
# List of feature pairs
feature_pairs = [
    ('Day_Mins', 'Day_Charge'),
    ('Eve_Mins', 'Eve_Charge'),
    ('Night_Mins', 'Night_Charge'),
    ('Intl_Mins', 'Intl_Charge'),
    # ...
]

# scatter plots for each pair
for x, y in feature_pairs:
    plt.figure(figsize=(8, 6))
    plt.scatter(df[x], df[y], alpha=0.5)
    plt.xlabel(x)
    plt.ylabel(y)
    plt.title(f'Scatter Plot of {x} vs {y}')
    plt.show()
```


Scatter Plot of Day_Mins vs Day_Charge




Scatter Plot of Eve_Mins vs Eve_Charge



Scatter Plot of Night_Mins vs Night_Charge





```
# pairplot
sns.set()
sns.pairplot(df[numeric_cols],size = 2 ,kind
='scatter',diag_kind='kde')
plt.show()
```

/opt/conda/lib/python3.10/site-packages/seaborn/axisgrid.py:2095:
UserWarning: The `size` parameter has been renamed to `height`; please
update your code.
warnings.warn(msg, UserWarning)

/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):

/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.

```
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
with pd.option_context('mode.use_inf_as_na', True):
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed
in a future version. Convert inf values to NaN before operating
instead.
```

instead.


```
with pd.option_context('mode.use_inf_as_na', True):  
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```

```
with pd.option_context('mode.use_inf_as_na', True):  
/opt/conda/lib/python3.10/site-packages/seaborn/_oldcore.py:1119:  
FutureWarning: use_inf_as_na option is deprecated and will be removed  
in a future version. Convert inf values to NaN before operating  
instead.
```


```
with pd.option_context('mode.use_inf_as_na', True):
```




```
# violin plot
for num_col in numeric_cols:
    sns.catplot(x='Churn', y=num_col, kind='violin', data=df,
palette='Set2')
    plt.title(f'Violin Plot of {num_col} by Churn')
    plt.xticks(rotation=45) # Rotate x labels if they overlap
    plt.show()
```





Violin Plot of Day_Mins by Churn







Violin Plot of Day_Charge by Churn












Violin Plot of Night_Mins by Churn





Violin Plot of Night_Calls by Churn




Violin Plot of Night_Charge by Churn







Violin Plot of Intl_Charge by Churn



Violin Plot of VMail_Message by Churn

